

# Pivotal Cloud Native Tour

---

- [Pivotal Cloud Native Tour](#)
- [Lab 1 - Getting Started with Spring Boot](#)
  - [Bootstrap the Project](#)
  - [Build and Run with Embedded Apache Tomcat](#)
  - [Refactor to Externalize the Config](#)
  - [Using Environment Variables for Config](#)
  - [Using Spring Profiles for Config](#)
  - [Resolving Conflicts](#)
- [Lab 2 - Introspection, Monitoring, and Metrics using Spring Boot Actuator](#)
  - [Set up the Actuator](#)
  - [Introspection Endpoints](#)
  - [Just read it - Build and Version Control Info](#)
  - [Health Indicators](#)
  - [Metrics](#)
- [Lab 3 - Composing Cloud Native Apps with Spring Cloud and Netflix OSS](#)
  - [Bootstrapping](#)
  - [Config Server](#)
  - [Eureka Server](#)
  - [Fortune Service](#)
  - [Fortune UI](#)
  - [Hystrix Dashboard](#)
- [Lab 4 - From Zero to Pushing Your First Application to Pivotal Cloud Foundry](#)
  - [Target](#)
  - [Build and Push!](#)
- [Demo - Lab 5 - Binding to Cloud Foundry Services](#)
  - [A Bit of Review](#)
  - [The Services Marketplace](#)
  - [Demo - Creating and Binding to a Service Instance](#)
- [Demo - Lab 6 - Monitoring Applications with Pivotal Cloud Foundry](#)
  - [Health](#)
- [Lab 7 - Push Cloud Native Apps to Pivotal Cloud Foundry](#)
  - [Build and Push!](#)

# Lab 1 - Getting Started with Spring Boot

---

## Bootstrap the Project

---

1. In your browser, visit <http://start.spring.io>.
2. Fill out the Project metadata fields as follows:

### Group

`io.pivotal.spring`

### Artifact

`hello-spring-boot`

3. In the Dependencies section, search for web, and choose **Web** from the autocomplete menu.
4. Click the Generate Project button. Your browser will download a zip file. Unpack that zip file at **\$COURSE\_HOME/labs/initial**.
5. Import the project's pom.xml into your editor/IDE of choice.
6. Add a `@RestController` annotation to the class `io.pivotal.spring.HelloSpringBootApplication`.
7. Add the following request handler to the class `io.pivotal.spring.HelloSpringBootApplication`:

```
@RequestMapping("/")
public String hello() {
    return "Hello World!";
}
```

## Build and Run with Embedded Apache Tomcat

---

Spring Boot will embed Apache Tomcat by default.

1. Build the application:

```
$ ./mvnw package
```

2. Run the application:

```
$ java -jar target/hello-spring-boot-0.0.1-SNAPSHOT.jar
```

You should see the application start up an embedded Apache Tomcat server on port 8080:

```
2015-12-06 17:17:30.507 INFO 60277 --- [           main] s.b.c.e.t.To
mcatEmbeddedServletContainer : Tomcat started on port(s): 8080 (http)
2015-12-06 17:17:30.511 INFO 60277 --- [           main] io.pivotal.s
pring.HelloSpringBootApplication : Started HelloSpringBootApplication
in 3.201 seconds (JVM running for 3.846)
```

3. Visit the application in the browser (<http://localhost:8080>), and you should see the following:

```
Hello World!
```

## Refactor to Externalize the Config

---

1. Rename `src/main/resources/application.properties` to `src/main/resources/application.yml`. Into that file, paste the following:

```
greeting: Hello
```

2. To the class `io.pivotal.spring.HelloSpringBootApplication`, add a `greeting` field and inject its value:

```
@Value("${greeting}")
String greeting;
```

3. Also `io.pivotal.spring.HelloSpringBootApplication`, change the return statement

of hello() to the following:

```
return String.format("%s World!", greeting);
```

4. Build the application:

```
$ ./mvnw package
```

5. Run the application:

```
$ java -jar target/hello-spring-boot-0.0.1-SNAPSHOT.jar
```

6. Visit the application in the browser (<http://localhost:8080>), and verify that the output is still the following:

```
Hello World!
```

7. Stop the application.

## Using Environment Variables for Config

---

1. Run the application again, this time setting the GREETING environment variable:

```
$ GREETING=Ohai java -jar target/hello-spring-boot-0.0.1-SNAPSHOT.jar
```

2. Visit the application in the browser (<http://localhost:8080>), and verify that the output has changed to the following:

```
Ohai World!
```

3. Stop the application.

## Using Spring Profiles for Config

---

1. Add a spanish profile to application.yml. Your finished configuration should reflect the following:

```
greeting: Hello

---

spring:
  profiles: spanish

greeting: Hola
```

2. Build the application:

```
$ ./mvnw package
```

3. Run the application, this time setting the SPRING\_PROFILES\_ACTIVE environment variable:

```
$ SPRING_PROFILES_ACTIVE=spanish java -jar target/hello-spring-boot-0.0.1-SNAPSHOT.jar
```

4. Visit the application in the browser (<http://localhost:8080>), and verify that the output has changed to the following:

```
Hola World!
```

5. Stop the application.

## Resolving Conflicts

---

1. Run the application, this time setting both the SPRING\_PROFILES\_ACTIVE and GREETING environment variables:

```
$ SPRING_PROFILES_ACTIVE=spanish GREETING=0hai java -jar target/hello-spring-boot-0.0.1-SNAPSHOT.jar
```

2. Visit the application in the browser (<http://localhost:8080>), and verify that the output has changed to the following:

```
Ohai World!
```

3. Visit <http://docs.spring.io/spring-boot/docs/current/reference/html/boot-features-external-config.html> to learn more about this outcome and the entire priority scheme for conflict resolution.

## Lab 2 - Introspection, Monitoring, and Metrics using Spring Boot Actuator

---

### Set up the Actuator

---

1. Continue with the project you completed in Lab 1.
2. To pom.xml add the following dependency to include the starter for Spring Boot Actuator:

```
<dependency>  
  <groupId>org.springframework.boot</groupId>  
  <artifactId>spring-boot-starter-actuator</artifactId>  
</dependency>
```

### Introspection Endpoints

---

1. Build the application:

```
$ ./mvnw package
```

2. Run the application:

```
$ java -jar target/hello-spring-boot-0.0.1-SNAPSHOT.jar
```

3. Try out the following endpoints. The output is omitted here because it can be quite large:

<http://localhost:8080/beans>

- Dumps all of the beans in the Spring context.

<http://localhost:8080/autoconfig>

- Dumps all of the auto-configuration performed as part of application bootstrapping.

<http://localhost:8080/configprops>

- Displays a collated list of all @ConfigurationProperties.

<http://localhost:8080/env>

- Dumps the application's shell environment as well as all Java system properties.

<http://localhost:8080/mappings>

- Dumps all URI request mappings and the controller methods to which they are mapped.

<http://localhost:8080/dump>

- Performs a thread dump.

<http://localhost:8080/trace>

- Displays trace information (by default the last few HTTP requests).

## Just read it - Build and Version Control Info

---

Spring Boot provides an endpoint (<http://localhost:8080/info>) that allows the exposure of arbitrary metadata.

One thing that it does well is expose information about the specific build and version control coordinates for a given deployment.

1. Add the following plugin to your Maven build. It will add Git branch and commit coordinates to the /info endpoint:

```
<plugin>
  <groupId>pl.project13.maven</groupId>
  <artifactId>git-commit-id-plugin</artifactId>
  <configuration>
    <dotGitDirectory>../../../../../.git</dotGitDirectory>
  </configuration>
</plugin>
```

## NOTE

The path ../../../../../.git refers to the .git directory at the root of the course materials.

2. Add the following properties to src/main/resources/application.yml:

```
info:
  build:
    artifact: @project.artifactId@
    name: @project.name@
    description: @project.description@
    version: @project.version@
```

These will add the project's Maven coordinates to the /info endpoint. The Spring Boot Maven plugin will cause them to automatically be replaced in the assembled JAR.

3. Build the application:

```
$ ./mvnw package
```

4. Run the application:

```
$ java -jar target/hello-spring-boot-0.0.1-SNAPSHOT.jar
```



5. Visit the application in the browser (<http://localhost:8080/info>), and verify that the output is similar to the following:

```
{
  "git" : {
    "commit" : {
      "time" : "2015-12-06T17:25:32-0500",
      "id" : "78b0e1b"
    },
    "branch" : "master"
  },
  "build" : {
    "version" : "0.0.1-SNAPSHOT",
    "name" : "hello-spring-boot",
    "artifact" : "hello-spring-boot",
    "description" : "Demo project for Spring Boot"
  }
}
```

## Health Indicators

---

Spring Boot provides an endpoint (<http://localhost:8080/health>) that allows for the notion of various health indicators.

1. Normally, when Spring Security is not enabled, the /health endpoint will only expose an UP or DOWN value. To simplify working with the endpoint for this lab, we will turn off its sensitivity. Add the following to `src/main/resources/application.yml`:

```
endpoints:
  health:
    sensitive: false
```

2. Create the class `io.pivotal.spring.FlappingHealthIndicator` and into it paste the following code:

```

@Component
public class FlappingHealthIndicator implements HealthIndicator{

    private Random random = new Random(System.currentTimeMillis());

    @Override
    public Health health() {
        int result = random.nextInt(100);
        if (result < 50) {
            return Health.down().withDetail("flapper", "failure").with
Detail("random", result).build();
        } else {
            return Health.up().withDetail("flapper", "ok").withDetail(
"random", result).build();
        }
    }
}

```

This demo health indicator will randomize the health check.

### 3. Build the application:

```
$ ./mvnw package
```

### 4. Run the application:

```
$ java -jar target/hello-spring-boot-0.0.1-SNAPSHOT.jar
```

### 5. Visit the application in the browser (<http://localhost:8080/health>), and verify that the output is similar to the following (and changes randomly!):

```
{
  status: "UP",
  flapping: {
    status: "UP",
    flapper: "ok",
    random: 69
  },
  diskSpace: {
    status: "UP",
    free: 113632186368,
    threshold: 10485760
  }
}
```

## Metrics

---

Spring Boot provides an endpoint (<http://localhost:8080/metrics>) that exposes several automatically collected metrics for your application. It also allows for the creation of custom metrics.

1. Create the class `io.pivotal.spring.GreetingService` and into it paste the following code:

```
@Component
public class GreetingService {

    @Autowired
    CounterService counterService;

    @Value("${greeting}")
    String greeting;

    public String getGreeting() {
        counterService.increment("counter.services.greeting.invoked");
        return greeting;
    }
}
```

This class is using the `@Autowired CounterService` to count the number of times that the `getGreeting()` method has been invoked.

## 2. Refactor the contents of the class

`io.pivotal.spring.HelloSpringBootApplication:`

```
@Autowired
private GreetingService greetingService;

@RequestMapping("/")
public String hello() {
    return String.format("%s World!", greetingService.getGreeting());
}

public static void main(String[] args) {
    SpringApplication.run(HelloSpringBootApplication.class, args);
}
```

`hello()` is now delegating the source of the greeting to our newly created `GreetingService`.

## 3. Build the application:

```
$ ./mvnw package
```

## 4. Run the application:

```
$ java -jar target/hello-spring-boot-0.0.1-SNAPSHOT.jar
```

5. Visit the application in the browser (<http://localhost:8080>) and refresh the page several times.
6. Now visit the `/metrics` endpoint (<http://localhost:8080/metrics>). Among the autogenerated metrics you should see a counter for the `GreetingService` invocations:

```
counter.services.greeting.invoked: 16,
```

To learn more about the autogenerated metrics, visit <http://docs.spring.io/spring-boot/docs/current/reference/html/production-ready-metrics.html>.

# Lab 3 - Composing Cloud Native Apps with Spring Cloud and Netflix OSS

---

In this lab we're going to build a Fortune Teller application from two microservices, one that serves up random Chinese fortune cookie fortunes, and one that presents a user interface. The resulting application looks like the following screenshot:

A screenshot of a web application interface. It features a large, bold, black title "Fortune Teller!" centered on a light gray background. Below the title, in a smaller, regular black font, is the text "Find out what the future holds...".

# Fortune Teller!

Find out what the future holds...

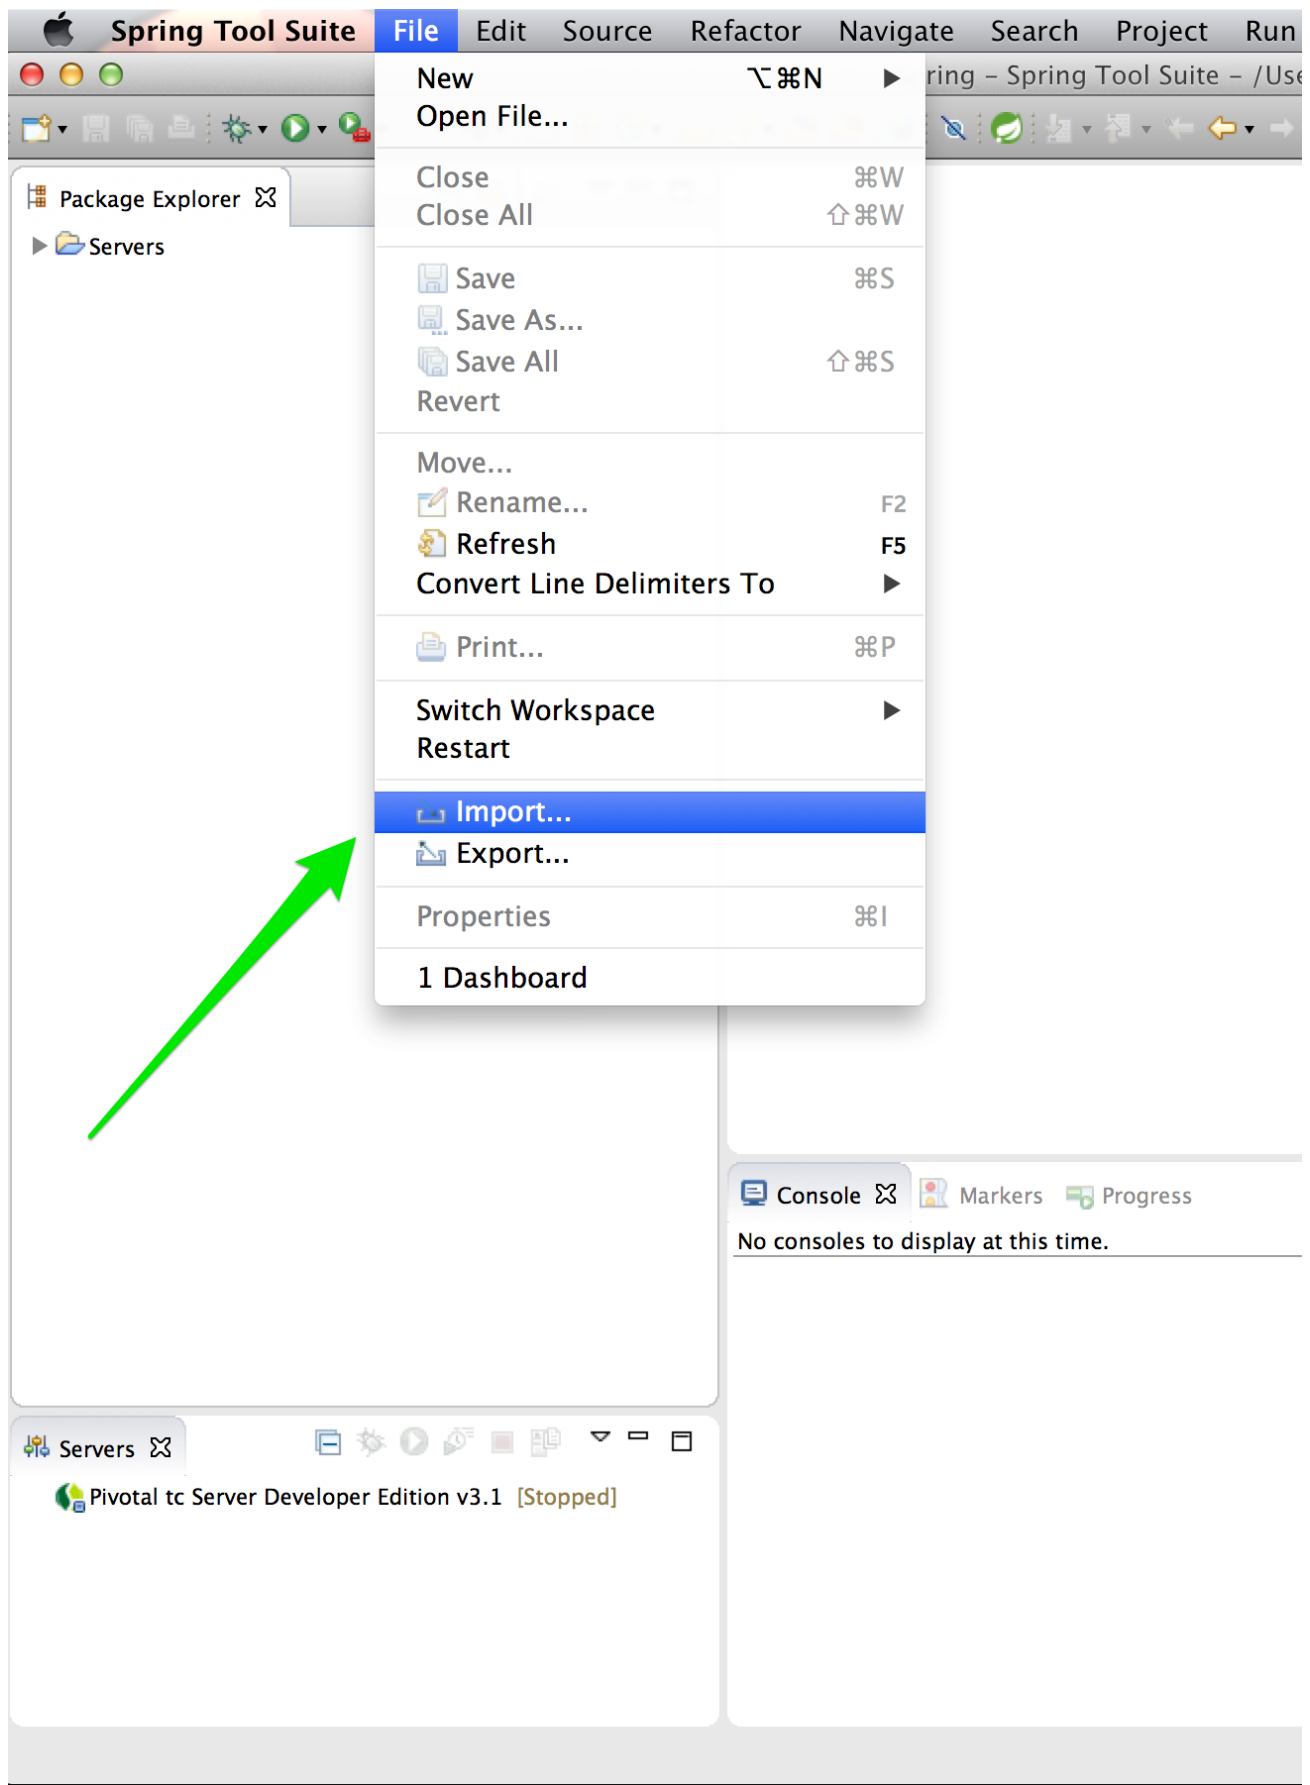
You cannot love life until you live the life you love.

We'll leverage libraries and services from Spring Cloud and Netflix OSS to help us compose the system.

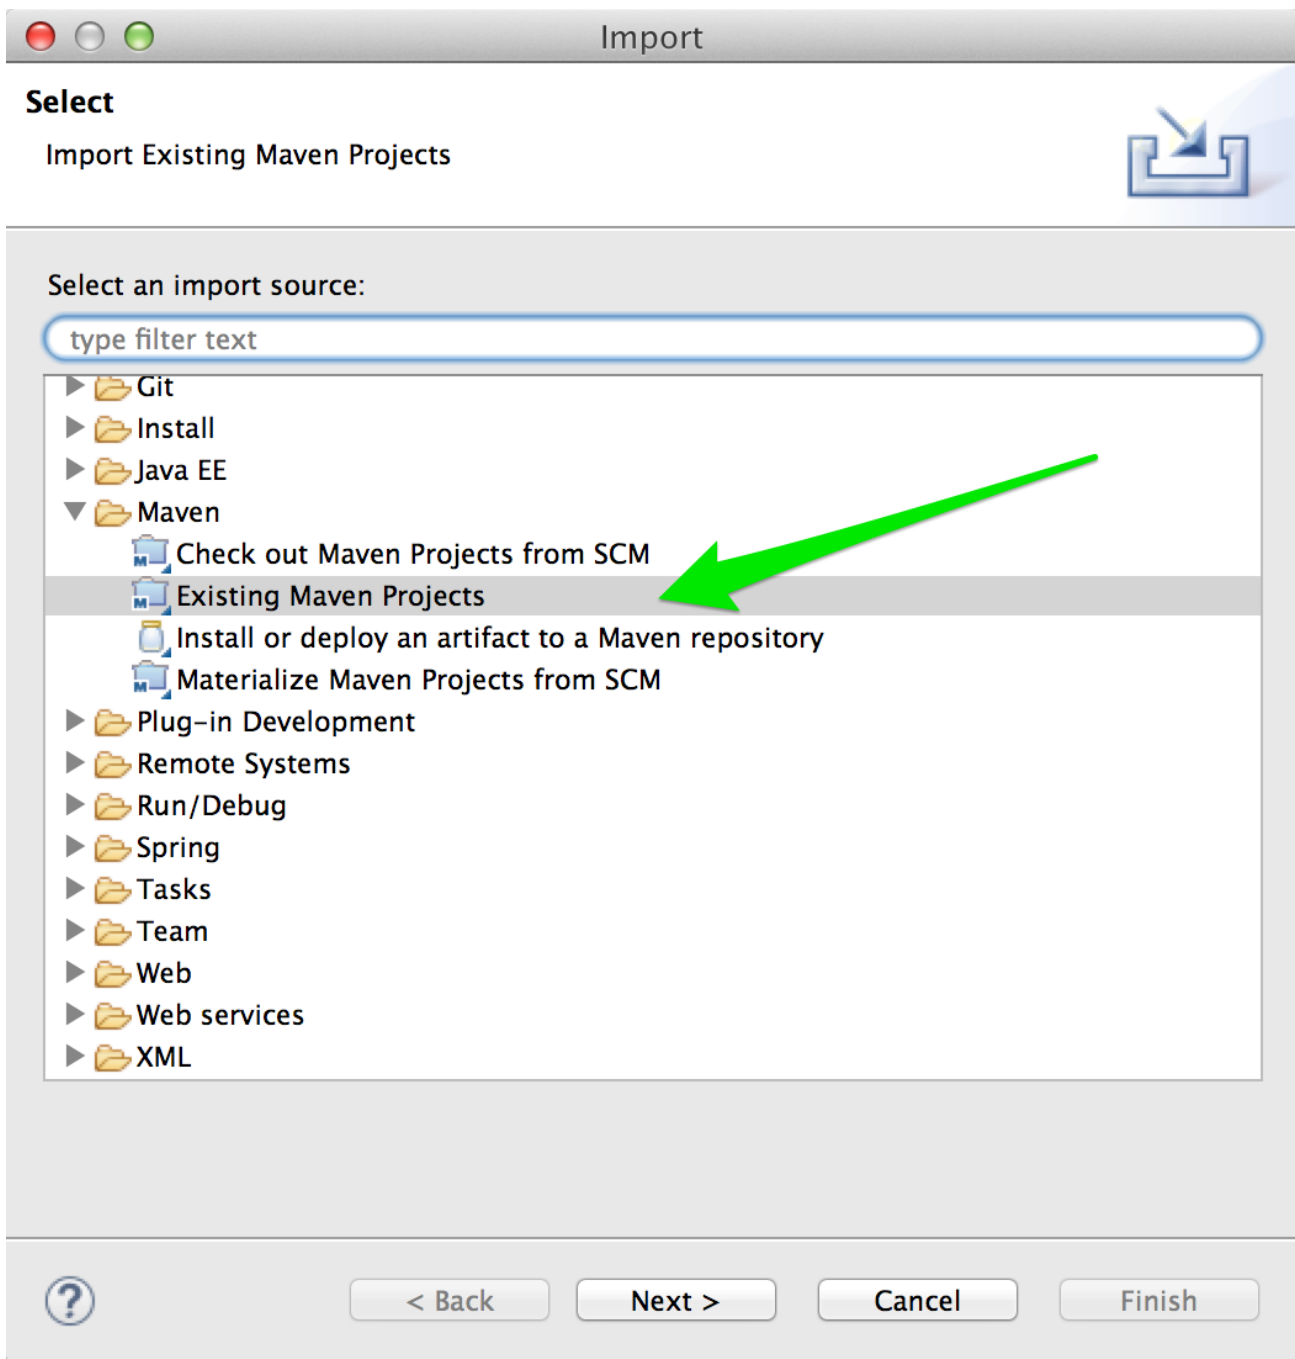
## Bootstrapping

---

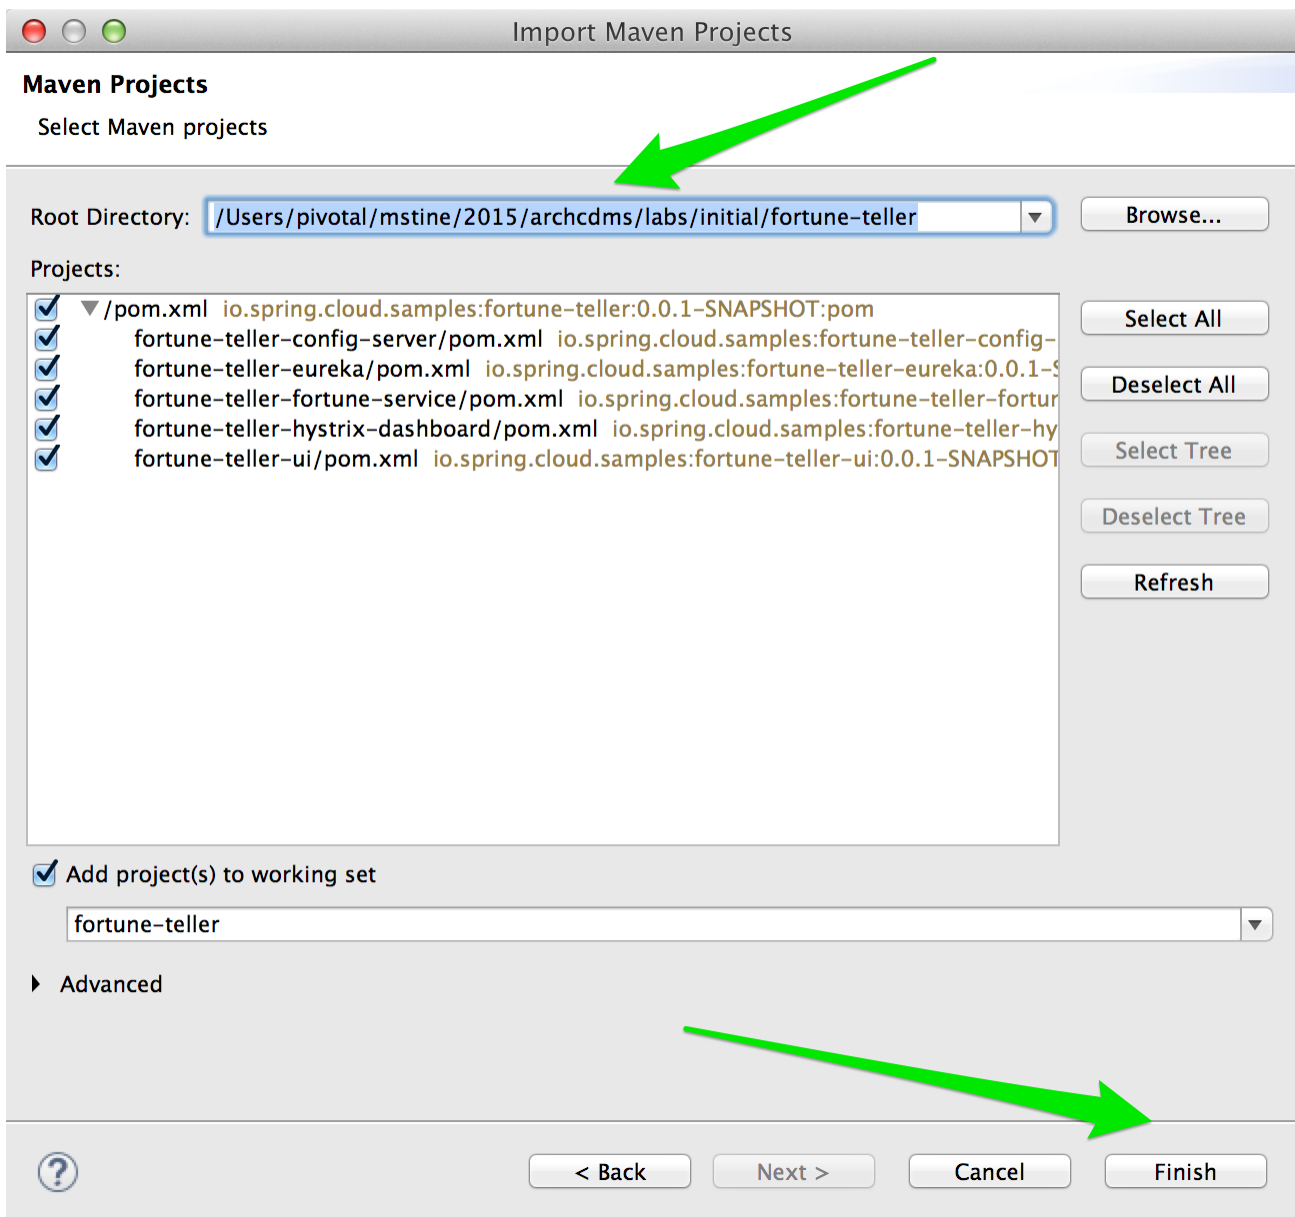
1. Choose File → Import:



2. Choose Maven → Existing Maven Projects:



3. Choose the project's pom.xml, found at `$COURSE_HOME/labs/initial/fortune-teller` and click Finish.



## Config Server

First we'll create a Spring Cloud Config Server to distribute shared configuration across our distributed system. This configuration will form the basis of how the entire system will compose.

1. In the fortune-teller-config-server module, add a `@EnableConfigServer` annotation to the class `io.spring.cloud.samples.fortuneteller.configserver.Application`.
2. Paste the following configuration properties into the file `src/main/resources/application.yml`:



```
server:
  port: 8888

spring:
  cloud:
    config:
      server:
        git:
          uri: https://github.com/mstine/config-repo.git
```

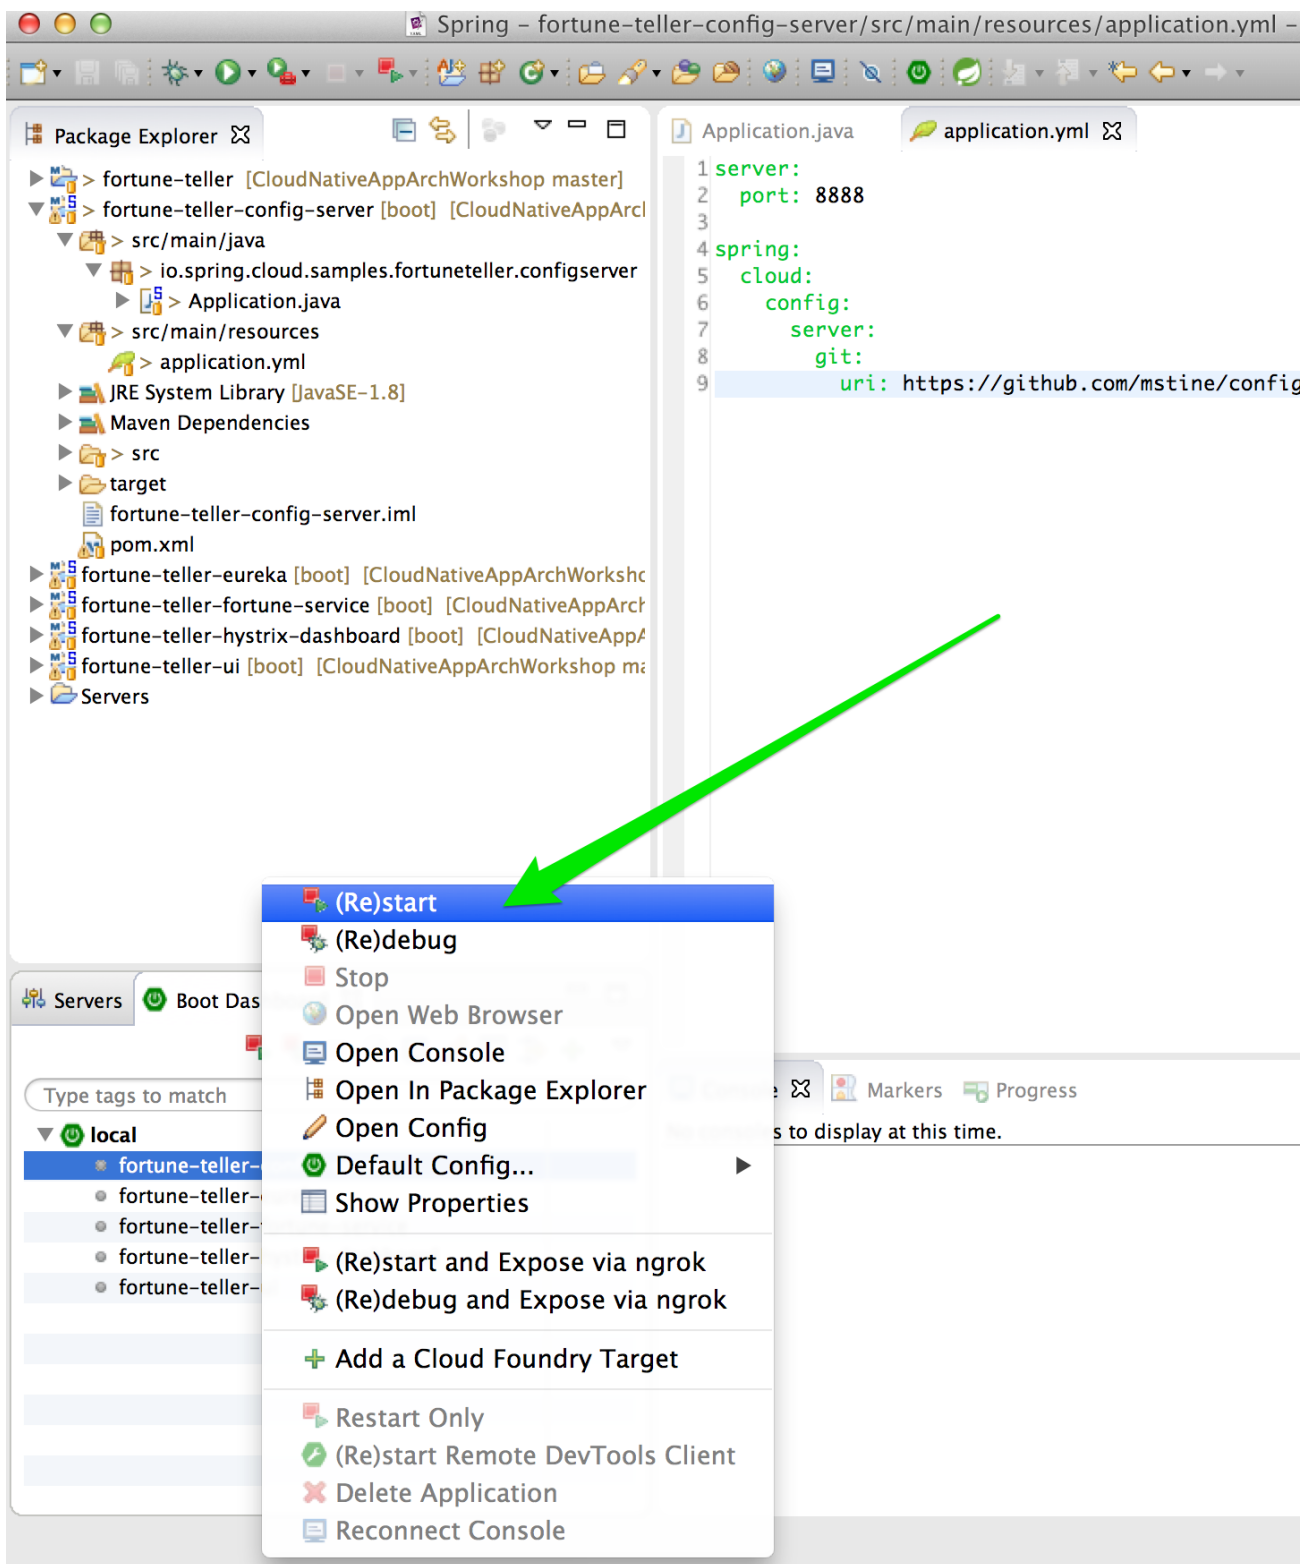
These properties tell the Config Server to listen on port 8888 and to source its configuration from the Git repository found at <https://github.com/mstine/config-repo>.

3. Browse the file

<https://github.com/mstine/config-repo/blob/master/application.yml#L1-L18>.

These lines tell all applications using the cloud Spring profile how to connect to and register themselves with Eureka.

4. In the Boot Dashboard, right click on fortune-teller-config-server and choose (Re)start:



5. Browse to <http://localhost:8888/application/default>. You should see output similar to the following:

config server output

6. Note that the response JSON defines a Spring PropertySource that contains Eureka configuration consistent with what you saw in the Git repository.

# Eureka Server

---

Next we'll create a Eureka service registry using Spring Cloud Netflix. This server will act as the address book our UI application will use to locate the Fortune Service.


1. In the fortune-teller-eureka module, add a `@EnableEurekaServer` annotation to the class  
`io.spring.cloud.samples.fortuneteller.eureka.Application`.
2. Paste the following configuration properties into the file  
`src/main/resources/application.yml`:

```
server:
  port: 8761

eureka:
  instance:
    hostname: localhost
  client:
    registerWithEureka: false
    fetchRegistry: false
    serviceUrl:
      defaultZone: http://${eureka.instance.hostname}:${server.port}/eureka/
```

These properties tell the Eureka Server to listen on port 8761 and to configure itself in what is essentially “standalone” mode.

3. As before, in the Boot Dashboard, right click on fortune-teller-eureka and choose (Re)start.
4. Browse to <http://localhost:8761>. You should see a UI similar to the following:

HOME    LAST 1000 SINCE STARTUP

### System Status

Environment	Current time	2015-02-18T18:49:44 -0800
Data center	Uptime	00:00
	Lease expiration enabled	false
	Renews threshold	0
	Renews (last min)	0

### DS Replicas

localhost

### Instances currently registered with Eureka

Application	AMIs	Availability Zones	Status
No instances available			

eureka 1

## Fortune Service

Now that our Spring Cloud/Netflix OSS infrastructure is bootstrapped, we'll begin development of our first microservice. The Fortune Service will serve up a list of all fortune cookie fortunes present in its database or a single random fortune.

1. In the fortune-teller-fortune-service module, create the package  
`io.spring.cloud.samples fortuneteller.fortuneservice.domain.`
2. Create a JPA domain class,  
`io.spring.cloud.samples fortuneteller.fortuneservice.domain.Fortune.`  
Into that class, paste the following code:

```

@Entity
@Table(name = "fortunes")
public class Fortune {

    @Id
    @GeneratedValue
    private Long id;

    @Column(nullable = false)
    private String text;

    public Long getId() {
        return id;
    }

    public void setId(Long id) {
        this.id = id;
    }

    public String getText() {
        return text;
    }

    public void setText(String text) {
        this.text = text;
    }
}

```

Add the appropriate imports using the IDE's quick fix feature.

This class represents the domain model for a fortune, having an unique identifier and the fortune's text.

3. Next, create the package

```
io.spring.cloud.samples.fortuneteller.fortuneservice.repositories.
```

4. Create a Spring Data JPA repository interface,  
 io.spring.cloud.samples.fortuneteller.fortuneservice.repositories.FortuneRepository.  
 Into that class, paste the following code:

```
public interface FortuneRepository extends PagingAndSortingRepository<
    Fortune, Long> {

    @Query("select fortune from Fortune fortune order by RAND()")
    public List<Fortune> randomFortunes(Pageable pageable);
}
```

5. Add the appropriate imports using the IDE's quick fix feature.

This interface will provide all of the basic boilerplate CRUD features we'll want for our Fortune domain class. The @Query annotation defines a SQL statement that will return a random ordering of the fortunes. We'll use this as the basis of returning a random fortune via our REST interface.

6. Next, create the package  
io.spring.cloud.samples.fortuneteller.fortuneservice.controllers.
7. Create a Spring MVC Controller class,  
io.spring.cloud.samples.fortuneteller.fortuneservice.controllers.FortuneController.  
Into that class, paste the following code:

```
@RestController
public class FortuneController {

    @Autowired
    FortuneRepository repository;

    @RequestMapping("/fortunes")
    public Iterable<Fortune> fortunes() {
        return repository.findAll();
    }

    @RequestMapping("/random")
    public Fortune randomFortune() {
        List<Fortune> randomFortunes = repository.randomFortunes(new PageRequest(0, 1));
        return randomFortunes.get(0);
    }
}
```

Add the appropriate imports using the IDE's quick fix feature.

This class provides the two REST endpoints of our microservice. The `randomFortune` method provides a `PageRequest` to limit the returned results to one.

8. Add `@EnableJpaRepositories` (to enable the Spring Data JPA repository we created) and `@EnableDiscoveryClient` (to switch on Eureka registration) annotations to `io.spring.cloud.samples.fortuneteller.fortuneservice.Application`.

Paste the following configuration properties into the file `src/main/resources/bootstrap.yml`:

```
spring:
  application:
    name: fortunes
  cloud:
    config:
      uri: ${vcap.services.config-service.credentials.uri:http://localhost:8888}
```

This configuration specifies the name of the application, which will be used to register the service in Eureka. It also tells the service how to locate the config server.

9. Paste the following configuration properties into the file `src/main/resources/application.yml`:

```
spring:
  jpa:
    hibernate:
      ddl-auto: create-drop
```

---

```
spring:
  profiles: default
  jpa:
    database: HSQL
```

---

```
spring:
  profiles: cloud
  jpa:
    database: POSTGRESQL
```

This configuration specifies the model to DDL setting to be used by Hibernate, as well as the Hibernate dialects to be used for local development and cloud.

10. As before, in the Boot Dashboard, right click on fortune-teller-fortune-service and choose (Re)start.

After the service has been running for ten seconds, refresh your Eureka browser tab. You should see a registration entry for the fortune service.

11. Visit <http://localhost:8080/random>. You should see output similar to the following:





localhost:8080/random

```
{  
  id: 1044,  
  text: "Let the deeds speak."  
}
```

## Fortune UI

---

Now that our Fortune microservice is running, we'll begin development of our second microservice. The Fortune UI will serve up a AngularJS single page application that consumes the fortune service.

1. In the fortune-teller-ui module, create the package  
`io.spring.cloud.samples fortuneteller.ui.services.fortunes`. This package will contain our integration with the fortune service.
2. Create the POJO  
`io.spring.cloud.samples.fortuneteller.ui.services.fortunes.Fortune`. Into that class, paste the following code:

```
public class Fortune {
    private Long id;
    private String text;

    public Fortune() {
    }

    public Fortune(Long id, String text) {
        this.id = id;
        this.text = text;
    }

    public Long getId() {
        return id;
    }

    public void setId(Long id) {
        this.id = id;
    }

    public String getText() {
        return text;
    }

    public void setText(String text) {
        this.text = text;
    }
}
```

We will use this class to unmarshal the response from the fortune service.

### 3. Create the Spring Bean

```
io.spring.cloud.samples fortuneteller.ui.services.fortunes.FortuneService.
```

Into that class, paste the following code:

```

@Service
public class FortuneService {

    @Autowired
    RestTemplate restTemplate;

    @HystrixCommand(fallbackMethod = "fallbackFortune")
    public Fortune randomFortune() {
        return restTemplate.getForObject("http://fortunes/random", Fortune.class);
    }

    private Fortune fallbackFortune() {
        return new Fortune(42L, "Your future is unclear.");
    }
}

```

Add the appropriate imports using the IDE's quick fix feature.

This class is our integration point with the fortune service. It uses a special `RestTemplate` bean that integrates with `Ribbon` from Netflix OSS. The argument passed to `getForObject`, `http://fortunes/random`, will be resolved by `Ribbon` to the actual address of the fortune service. This method is also protected by a `Hystrix` circuit breaker using the `@HystrixCommand` annotation. If the circuit is tripped to open (due to an unhealthy fortune service), the `fallbackFortune` method will return a dummy response.

4. Next, create the package

```
io.spring.cloud.samples.fortuneteller.ui.controllers.
```

5. Create a Spring MVC Controller class,

```
io.spring.cloud.samples.fortuneteller.ui.controllers.UiController.
```

Into that class, paste the following code:

```

@RestController
public class UiController {

    @Autowired
    FortuneService service;

    @RequestMapping("/random")
    public Fortune randomFortune() {
        return service.randomFortune();
    }
}

```

6. Add the appropriate imports using the IDE's quick fix feature.

This class provides the REST endpoint that will be consumed by our AngularJS UI.

7. Add `@EnableCircuitBreaker` (to switch on Hystrix circuit breakers) and `@EnableDiscoveryClient` (to switch on Eureka registration) annotations to `io.spring.cloud.samples.fortuneteller.ui.Application`.
8. Paste the following configuration properties into the file `src/main/resources/bootstrap.yml`:

```

spring:
  application:
    name: ui
  cloud:
    config:
      uri: ${vcap.services.config-service.credentials.uri:http://local
host:8888}

```

This configuration specifies the name of the application, which will be used to register the service in Eureka. It also tells the service how to locate the config server.

9. Paste the following configuration properties into the file `src/main/resources/application.yml`:

```
server:
  port: 8081

eureka:
  client:
    registryFetchIntervalSeconds: 5

fortunes:
  ribbon:
    ServerListRefreshInterval: 5000
```

This configuration instructs the application to listen on port 8081. It also provides low cache refresh settings for the Eureka client and Ribbon to make the demo more responsive.

10. As before, in the Boot Dashboard, right click on fortune-teller-ui and choose (Re)start.

After the service has been running for ten seconds, refresh your Eureka browser tab. You should see a registration entry for the Fortune UI.

11. Visit <http://localhost:8081>. You should see output similar to the following:

# Fortune Teller!

Find out what the future holds...

You cannot love life until you live the life you love.

# Hystrix Dashboard

---

Finally, we'll add one more piece of infrastructure — a dashboard that allows us to monitor our circuit breakers — using Spring Cloud Netflix.

1. In the fortune-teller-hystrix-dashboard module, add `@EnableHystrixDashboard` and `@Controller` annotations to the class `io.spring.cloud.samples.fortuneteller.hystrixdashboard.Application`.
2. Add the following request handler method to `io.spring.cloud.samples.fortuneteller.hystrixdashboard.Application`:

```
@RequestMapping("/")
public String home() {
    return "forward:/hystrix";
}
```

3. Add the appropriate imports using the IDE's quick fix feature.

This request handler will allow us to hit the default endpoint our application to interact with the Hystrix Dashboard.

4. Paste the following configuration properties into the file `src/main/resources/application.yml`:

```
server:
  port: 7979
```

This configuration instructs the application to listen on port 7979.

5. As before, in the Boot Dashboard, right click on fortune-teller-hystrix-dashboard and choose (Re)start.
6. Visit <http://localhost:7979>. You should see output similar to the following:



## Hystrix Dashboard

*Cluster via Turbine (default cluster):* `http://turbine-hostname:port/turbine.stream`

*Cluster via Turbine (custom cluster):* `http://turbine-hostname:port/turbine.stream?cluster=[clusterName]`

*Single Hystrix App:* `http://hystrix-app:port/hystrix.stream`

Delay:  ms

Title:

7. Paste <http://localhost:8081/hystrix.stream> into the text field and click Monitor Stream. You should see output similar to the following:

Hystrix Stream: <http://localhost:8081/hystrix.stream>



**Circuit** Sort: [Error then Volume](#) | [Alphabetical](#) | [Volume](#) | [Error](#) | [Mean](#) | [Median](#) | [90](#) | [99](#) | [99.5](#) | [Success](#) | [Short-Circuited](#) | [Timeout](#) | [Rejected](#) | [Failure](#) | [Error %](#)

randomFortune  
0 0 0.0 %  
0 0 0  
Host: 0.0/s  
Cluster: 0.0/s  
Circuit Closed  
90th 0ms  
99th 0ms  
99.5th 0ms  
Hosts 1  
Median 0ms  
Mean 0ms

**Thread Pools** Sort: [Alphabetical](#) | [Volume](#) |

FortuneService  
Host: 0.0/s  
Cluster: 0.0/s  
Active 0  
Queued 0  
Pool Size 3  
Max Active 0  
Executions 0  
Queue Size 5

# Lab 4 - From Zero to Pushing Your First Application to Pivotal Cloud Foundry

# Target

---

1. If you haven't already, download the latest release of the Cloud Foundry CLI from <https://github.com/cloudfoundry/cli/releases> for your operating system and install it.
2. Set the API target for the CLI:

```
$ cf api api.run.pivotal.io
```

3. Login to Pivotal Web Services:

```
$ cf login
```

Follow the prompts.

## Build and Push!

---

Change to the Spring Music sample application directory:

```
$ cd $COURSE_HOME/labs/complete/spring-music
```

Using the Gradle Wrapper, build and package the application:

```
$ ./gradlew assemble
```

The Gradle Wrapper will automatically download the appropriate version of Gradle for this project along with all of Spring Music's dependencies. This may take a few moments.

3. Push the application!

```
$ cf push
```

You should see output similar to the following listing. Take a look at the listing



callouts for a play-by-play of what's happening:

```
Using manifest file
/Users/pivotal/mstine/2015/0Reilly/CloudNativeArchitectureClass/day_01
/session_02/lab_05/spring-music/manifest.yml (1)
```

```
Creating app spring-music in org oreilly-class / space instructor as m
stine@pivotal.io...
OK (2)
```

```
Creating route spring-music-hippest-shaman.cfapps.io...
OK (3)
```

```
Binding spring-music-hippest-shaman.cfapps.io to spring-music...
OK (4)
```

```
Uploading spring-music... (5)
Uploading app files from: /Users/pivotal/mstine/2015/0Reilly/CloudNati
veArchitectureClass/day_01/session_02/lab_05/spring-music/build/libs/s
pring-music.war
Uploading 569.7K, 90 files
Done uploading
OK
```

```
Starting app spring-music in org oreilly-class / space instructor as m
stine@pivotal.io... (6)
-----> Downloaded app package (21M)
-----> Java Buildpack Version: v2.6.1 | https://github.com/cloudfoundry/java-buildpack.git#2d92e70
-----> Downloading Open Jdk JRE 1.8.0_31 from https://download.run.pivotal.io/openjdk/lucid/x86\_64/openjdk-1.8.0\_31.tar.gz (1.3s)
      Expanding Open Jdk JRE to .java-buildpack/open_jdk_jre (1.2s) (
7)
-----> Downloading Spring Auto Reconfiguration 1.7.0_RELEASE from https://download.run.pivotal.io/auto-reconfiguration/auto-reconfiguration-1.7.0\_RELEASE.jar (0.1s)
-----> Downloading Tomcat Instance 8.0.18 from https://download.run.pivotal.io/tomcat/tomcat-8.0.18.tar.gz (0.4s)
      Expanding Tomcat to .java-buildpack/tomcat (0.1s) (8)
-----> Downloading Tomcat Lifecycle Support 2.4.0_RELEASE from https://download.run.pivotal.io/tomcat-lifecycle-support/tomcat-lifecycle-support-2.4.0\_RELEASE.jar (0.0s)
-----> Downloading Tomcat Logging Support 2.4.0_RELEASE from https://download.run.pivotal.io/tomcat-logging-support/tomcat-logging-support-2.4.0\_RELEASE.jar (0.0s)
-----> Downloading Tomcat Access Logging Support 2.4.0_RELEASE from https://download.run.pivotal.io/tomcat-access-logging-support/tomcat-access-logging-support-2.4.0\_RELEASE.jar (0.0s)
```

```
ess-logging-support-2.4.0_RELEASE.jar (0.0s)
-----> Uploading droplet (66M) (9)
```

```
0 of 1 instances running, 1 starting
1 of 1 instances running
```

App started

OK

App spring-music was started using this command `JAVA\_HOME=\$PWD/.java-buildpack/open\_jdk\_jre JAVA\_OPTS="-Djava.io.tmpdir=\$TMPDIR -XX:OnOutOfMemoryError=\$PWD/.java-buildpack/open\_jdk\_jre/bin/killjava.sh -Xmx382293K -Xms382293K -XX:MaxMetaspaceSize=64M -XX:MetaspaceSize=64M -Xss995K -Daccess.logging.enabled=false -Dhttp.port=\$PORT" \$PWD/.java-buildpack/tomcat/bin/catalina.sh run` (10)

Showing health and status for app spring-music in org oreilly-class / space instructor as mstine@pivotal.io... (11)

OK

```
requested state: started
instances: 1/1
usage: 512M x 1 instances
urls: spring-music-hippest-shaman.cfapps.io
last uploaded: Fri Feb 13 15:43:08 UTC 2015
```

	state	since	cpu	memory	disk
#0	running	2015-02-13 09:43:55 AM	0.0%	394.5M of 512M	131.1M of 1G

1. The CLI is using a manifest to provide necessary configuration details such as application name, memory to be allocated, and path to the application artifact. Take a look at manifest.yml to see how.
2. In most cases, the CLI indicates each Cloud Foundry API call as it happens. In this case, the CLI has created an application record for Spring Music in your assigned space.
3. All HTTP/HTTPS requests to applications will flow through Cloud Foundry's front-end router called (Go)Router. Here the CLI is creating a route with random word tokens inserted (again, see manifest.yml for a hint!) to prevent route collisions across the default cfapps.io domain.

4. Now the CLI is binding the created route to the application. Routes can actually be bound to multiple applications to support techniques such as blue-green deployments.
5. The CLI finally uploads the application bits to PWS. Notice that it's uploading 90 files! This is because Cloud Foundry actually explodes a ZIP artifact before uploading it for caching purposes.
6. Now we begin the staging process. The Java Buildpack is responsible for assembling the runtime components necessary to run the application.
7. Here we see the version of the JRE that has been chosen and installed.
8. And here we see the version of Tomcat that has been chosen and installed.
9. The complete package of your application and all of its necessary runtime components is called a droplet. Here the droplet is being uploaded to PWS's internal blobstore so that it can be easily copied to one or more Droplet Execution Agents (DEA's) for execution.
10. The CLI tells you exactly what command and argument set was used to start your application.
11. Finally the CLI reports the current status of your application's health. You can get the same output at any time by typing `cf app spring-music`.
12. Visit the application in your browser by hitting the route that was generated by the CLI:

Spring Music 🎵

ⓘ

## Albums

[ view as: | sort by: [title](#) [artist](#) [year](#) [genre](#) ↕ | [+add an album](#) ]

Profiles: cloud,in-memory  
Services:

<div>IV</div> <div>Led Zeppelin</div> <div>1971</div> <div>Rock</div> <div>⚙</div>	<div>Nevermind</div> <div>Nirvana</div> <div>1991</div> <div>Rock</div> <div>⚙</div>	<div>What's Going On</div> <div>Marvin Gaye</div> <div>1971</div> <div>Rock</div> <div>⚙</div>	<div>Are You Experienced?</div> <div>Jimi Hendrix Experience</div> <div>1967</div> <div>Rock</div> <div>⚙</div>
<div>The Joshua Tree</div> <div>U2</div> <div>1987</div> <div>Rock</div> <div>⚙</div>	<div>Abbey Road</div> <div>The Beatles</div> <div>1969</div> <div>Rock</div> <div>⚙</div>	<div>Rumours</div> <div>Fleetwood Mac</div> <div>1977</div> <div>Rock</div> <div>⚙</div>	<div>Sun Sessions</div> <div>Elvis Presley</div> <div>1976</div> <div>Rock</div> <div>⚙</div>
<div>Thriller</div> <div>Michael Jackson</div> <div>1982</div> <div>Pop</div> <div>⚙</div>	<div>Exile on Main Street</div> <div>The Rolling Stones</div> <div>1972</div> <div>Rock</div> <div>⚙</div>	<div>Born to Run</div> <div>Bruce Springsteen</div> <div>1975</div> <div>Rock</div> <div>⚙</div>	<div>London Calling</div> <div>The Clash</div> <div>1980</div> <div>Rock</div> <div>⚙</div>

Be sure to click on the “information icon” in the top right-hand corner of the UI. This

gives you important information about the state of the currently running Spring Music instance, including what Spring Profiles are turned on and what Cloud Foundry services are bound. It will become important in the next lab!

## Demo - Lab 5 - Binding to Cloud Foundry Services

---

The Spring Music application was designed to illustrate the ease with which various types of data services can be bound to and utilized by Spring applications running on Cloud Foundry. In this lab, we'll be binding the application to both PostgreSQL and MongoDB databases.

Cloud Foundry services are managed through two primary types of operations:

### **Create/Delete**

These operations create or delete instances of a service. For a database this could mean creating/deleting a schema in an existing multitenant cluster or creating/deleting a dedicated database cluster.

### **Bind/Unbind**

These operations create or delete unique credential sets for an existing service instance that can then be injected into the environment of an application instance.

## A Bit of Review

---

Your instance of Spring Music should still be running from the end of Lab 7. Visit the application in your browser by hitting the route that was generated by the CLI:

Spring Music 🎵

Albums

[ view as: [grid icon] [list icon] | sort by: title artist year genre ^ | +add an album ]

Profiles: cloud,in-memory  
Services:

<div>IV</div> <div>Led Zeppelin</div> <div>1971</div> <div>Rock</div> <div>⚙</div>	<div>Nevermind</div> <div>Nirvana</div> <div>1991</div> <div>Rock</div> <div>⚙</div>	<div>What's Going On</div> <div>Marvin Gaye</div> <div>1971</div> <div>Rock</div> <div>⚙</div>	<div>Are You Experienced?</div> <div>Jimi Hendrix Experience</div> <div>1967</div> <div>Rock</div> <div>⚙</div>
<div>The Joshua Tree</div> <div>U2</div> <div>1987</div> <div>Rock</div> <div>⚙</div>	<div>Abbey Road</div> <div>The Beatles</div> <div>1969</div> <div>Rock</div> <div>⚙</div>	<div>Rumours</div> <div>Fleetwood Mac</div> <div>1977</div> <div>Rock</div> <div>⚙</div>	<div>Sun Sessions</div> <div>Elvis Presley</div> <div>1976</div> <div>Rock</div> <div>⚙</div>
<div>Thriller</div> <div>Michael Jackson</div> <div>1982</div> <div>Pop</div> <div>⚙</div>	<div>Exile on Main Street</div> <div>The Rolling Stones</div> <div>1972</div> <div>Rock</div> <div>⚙</div>	<div>Born to Run</div> <div>Bruce Springsteen</div> <div>1975</div> <div>Rock</div> <div>⚙</div>	<div>London Calling</div> <div>The Clash</div> <div>1980</div> <div>Rock</div> <div>⚙</div>

The information dialog in the top right-hand corner indicates that we’re currently running with an in-memory database, and that we’re not bound to any services. Let’s change that.

# The Services Marketplace

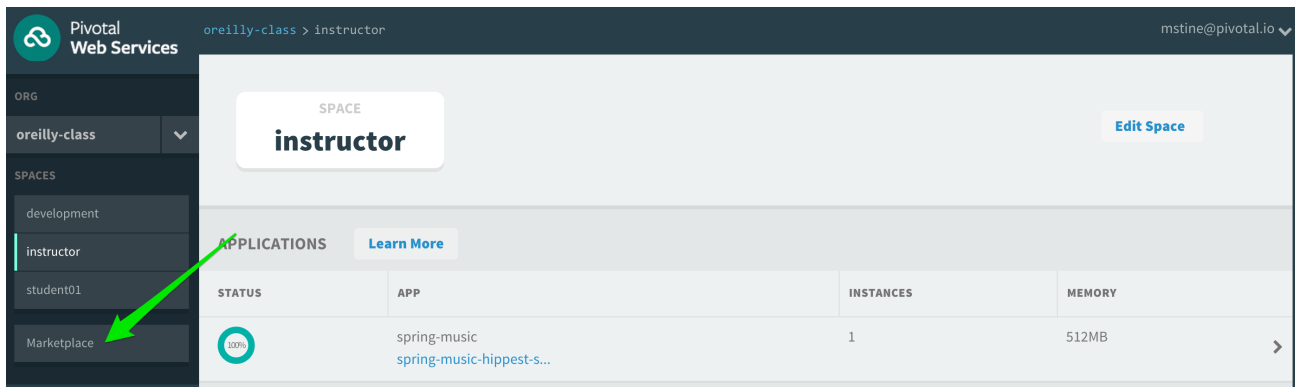
There are two ways to discover what services are available on Pivotal Web Services. The first is available on any instance of Cloud Foundry: the CLI. Just type:

```
$ cf marketplace
```

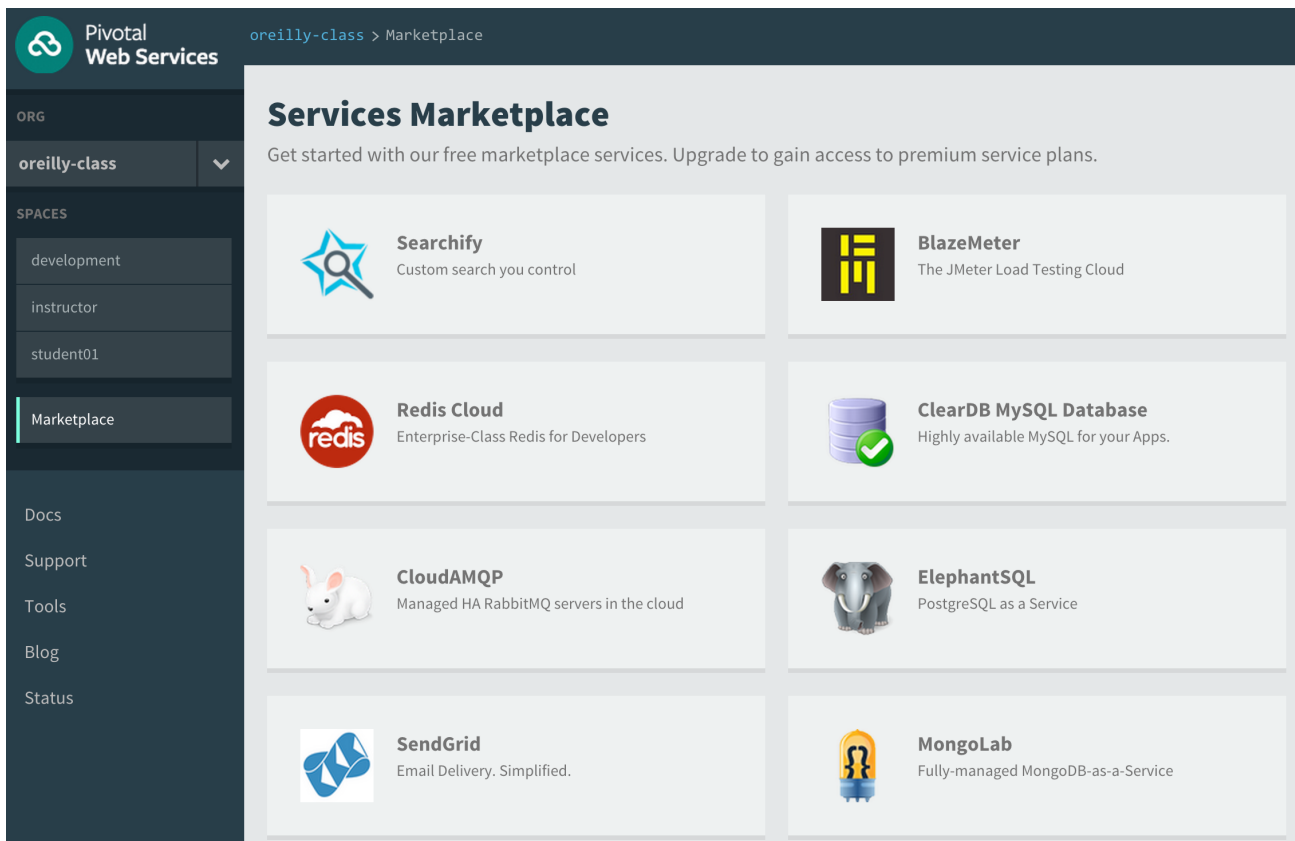
and you’ll get a list of services, their available plans, and descriptions. On Pivotal Web Services, the “free” tier of plans is normally the first one listed (e.g. turtle for elephantsql).

The second way is specific to PWS’s Application Manager UI. If you haven’t already, login to it by visiting <http://console.run.pivotal.io>.

Click on the “Marketplace” link:



and you'll see the same service/plan/description listing in the browser:



## Demo - Creating and Binding to a Service Instance

1. Let's begin by creating a PostgreSQL instance provided by ElephantSQL. From the CLI, let's create a free ElephantSQL service instance:

```
$ cf create-service elephantsql turtle spring-music-db
Creating service spring-music-db in org oreilly-class / space instructor as mstine@pivotal.io...
OK
```

2. Next we'll bind the newly created instance to our spring-music application:

```
$ cf bs spring-music spring-music-db
Binding service spring-music-db to app spring-music in org oreilly-class / space instructor as mstine@pivotal.io...
OK
TIP: Use 'cf restage' to ensure your env variable changes take effect
```

3. Notice the admonition to Use 'cf restage' to ensure your env variable changes take effect. Let's take a look at the environment variables for our application to see what's been done. We can do this by typing:

```
$ cf env spring-music
```

The subset of the output we're interested in is located near the very top, titled System-Provided:



```
System-Provided:
{
  "VCAP_SERVICES": { (1)
    "elephantsql": [
      {
        "credentials": {
          "max_conns": "5",
          "uri": "postgres://xnxzdeao:ntLT0M10gz014oXnKkkjeIqLn3Td2EgH@baba
r.elephantsql.com:5432/xnxzdeao" (2)
        },
        "label": "elephantsql",
        "name": "spring-music-db",
        "plan": "turtle",
        "tags": [
          "Data Stores",
          "Cloud Databases",
          "Developer Tools",
          "Data Store",
          "postgresql",
          "relational",
          "New Product"
        ]
      }
    ]
  }
}
```

VCAP\_SERVICES is a special Cloud Foundry environment variable that contains a JSON document containing all of the information for any services bound to an application.

Notice here the unique URI for this instance of PostgreSQL that spring-music has been bound to.

4. Now let's restage the application, which cycles our application back through the staging/buildpack process before redeploying the application.[1]

```
$ cf restage spring-music
```

Once the application is running again, revisit or refresh the browser tab where you have the Spring Music application loaded:

Spring Music

Albums

[ view as: | sort by: title artist year genre | +add an album ]

IV  
Led Zeppelin  
1971  
Rock  
⚙

Nevermind  
Nirvana  
1991  
Rock  
⚙

What's Going On  
Marvin Gaye  
1971  
Rock  
⚙

Are You Experienced?  
Jimi Hendrix Experience  
1967  
Rock  
⚙

The Joshua Tree  
U2  
1987  
Rock  
⚙

Abbey Road  
The Beatles  
1969  
Rock  
⚙

Rumours  
Fleetwood Mac  
1977  
Rock  
⚙

Sun Sessions  
Elvis Presley  
1976  
Rock  
⚙

Thriller  
Michael Jackson  
1982  
Pop  
⚙

Exile on Main Street  
The Rolling Stones  
1972  
Rock  
⚙

Born to Run  
Bruce Springsteen  
1975  
Rock  
⚙

London Calling  
The Clash  
1980  
Rock  
⚙

Profiles: cloud,postgres,postgres-cloud

Services: spring-music-db

As you can see from the information dialog, the application is now utilizing a PostgreSQL database via the spring-music-db service.

5. Let’s take a direct look at the data in PostgreSQL by utilizing the Manage link in PWS App Manager:

100%

spring-music

1

512MB

>

SERVICES

Add Service

SERVICE INSTANCE	SERVICE PLAN	BOUND APPS
<div>spring-music-db</div> <div><a href="#">Manage</a>   <a href="#">Documentation</a>   <a href="#">Support</a>   <a href="#">Delete</a></div>	ElephantSQL Tiny Turtle	1

6. Click on the Browser link:

ElephantSQL Console

b285234c-2a10-428d-b5d6-b3a04b76faa7

Management

Log

Browser

Backup

Database info

Reset database

URL	postgres://mzdlkpex:DxA4wTcZbYuNuym0lOou8qheJmijQkE@pellefant.db.elephantsql.com:5432/mzdlkpex	
Hostname	pellefant-01.db.elephantsql.com	
Database name	mzdlkpex	Use this for Maintenance DB in PgAdmin
Username	mzdlkpex	
Password	DxA4wTcZbYuNuym0lOou8qheJmijQkE	

7. Paste in the query `select * from album;` and click Execute:

## ElephantSQL Console b285234c-2a10-428d-b5d6-b3a04b76faa7

[Management](#) [Log](#) [Browser](#) [Backup](#)

### SQL Browser

select \* from album;

Execute

Tables

id	albumid	artist	genre	releaseyear	title	trackcount
81666f54-7cc3-46fb-8225-d098271db816		Nirvana	Rock	1991	Nevermind	0
a00b077d-c187-4b89-9581-0de989748659		The Beach Boys	Rock	1966	Pet Sounds	0
6b2f6cba-1cb5-4ea8-8d7d-089801cf0704		Marvin Gaye	Rock	1971	What's Going On	0
061e9cd1-5930-4aef-87f6-21f3d6bde2b2		Jimi Hendrix Experience	Rock	1967	Are You Experienced?	0

# Demo - Lab 6 - Monitoring Applications with Pivotal Cloud Foundry

Cloud Foundry provides several built-in mechanisms that allow us to monitor our applications' state changes and behavior. Additionally, Cloud Foundry actively monitors the health of our application processes and will restart them should they crash. In this lab, we'll explore a few of these mechanisms.

## Events

Cloud Foundry only allows application configuration to be modified via its API. This gives application operators confidence that all changes to application configuration are known and auditable. It also reduces the number of causes that must be considered when problems arise.

All application configuration changes are recorded as events. These events can be viewed via the Cloud Foundry API, and viewing is facilitated via the CLI.

Take a look at the events that have transpired so far for our deployment of `cf-scale-boot`:

```
$ cf events cf-scale-boot
```

Getting events for app cf-scale-boot in org oreilly-class / space instructor as [mstine@pivotal.io](mailto:mstine@pivotal.io)...

time event actor description

2015-02-13T15:18:33.00-0600 audit.app.update [mstine@pivotal.io](mailto:mstine@pivotal.io) instances: 1 (6)

2015-02-13T15:04:34.00-0600 audit.app.update [mstine@pivotal.io](mailto:mstine@pivotal.io) instances: 5 (5)

2015-02-13T12:56:35.00-0600 audit.app.update [mstine@pivotal.io](mailto:mstine@pivotal.io) state: STARTED (4)

2015-02-13T12:56:26.00-0600 audit.app.update [mstine@pivotal.io](mailto:mstine@pivotal.io) (3)

2015-02-13T12:56:26.00-0600 audit.app.map-route [mstine@pivotal.io](mailto:mstine@pivotal.io) (2)

2015-02-13T12:56:24.00-0600 audit.app.create [mstine@pivotal.io](mailto:mstine@pivotal.io) instances: 1, memory: 512, state: STOPPED, environment\_json: PRIVATE DATA HIDDEN (1)

Events are sorted newest to oldest, so we'll start from the bottom. Here we see the app.create event, which created our application's record and stored all of its metadata (e.g. memory: 512).

The app.map-route event records the incoming request to assign a route to our application.

This app.update event records the resulting change to our applications metadata.

This app.update event records the change of our application's state to STARTED.

Remember scaling the application up? This app.update event records the metadata change instances: 5.

And here's the app.update event recording our scaling of the application back down with instances: 1.

Let's explicitly ask for the application to be stopped:

```
$ cf stop cf-scale-boot
```

Stopping app cf-scale-boot in org oreilly-class / space instructor as [mstine@pivotal.io](mailto:mstine@pivotal.io)...

OK

Now, examine the additional app.update event:

```
$ cf events cf-scale-boot
```

Getting events for app cf-scale-boot in org oreilly-class / space instructor as [mstine@pivotal.io](mailto:mstine@pivotal.io)...

time event actor description

2015-02-13T15:59:10.00-0600 audit.app.update [mstine@pivotal.io](#) state: STOPPED

2015-02-13T15:18:33.00-0600 audit.app.update [mstine@pivotal.io](#) instances: 1

2015-02-13T15:04:34.00-0600 audit.app.update [mstine@pivotal.io](#) instances: 5

2015-02-13T12:56:35.00-0600 audit.app.update [mstine@pivotal.io](#) state: STARTED

2015-02-13T12:56:26.00-0600 audit.app.update [mstine@pivotal.io](#)

2015-02-13T12:56:26.00-0600 audit.app.map-route [mstine@pivotal.io](#)

2015-02-13T12:56:24.00-0600 audit.app.create [mstine@pivotal.io](#) instances: 1,  
memory: 512, state: STOPPED, environment\_json: PRIVATE DATA HIDDEN

Start the application again:

```
$ cf start cf-scale-boot
```

Starting app cf-scale-boot in org oreilly-class / space instructor as  
[mstine@pivotal.io](#)...

0 of 1 instances running, 1 starting

0 of 1 instances running, 1 starting

0 of 1 instances running, 1 starting

1 of 1 instances running

App started

OK

App cf-scale-boot was started using this command `JAVA_HOME=$PWD/.java-buildpack/open_jdk_jre JAVA_OPTS="-Djava.io.tmpdir=$TMPDIR -XX:OnOutOfMemoryError=$PWD/.java-buildpack/open_jdk_jre/bin/killjava.sh -Xmx382293K -Xms382293K -XX:MaxMetaspaceSize=64M -XX:MetaspaceSize=64M -Xss995K" SERVER_PORT=$PORT $PWD/.java-buildpack/spring_boot_cli/bin/spring run app.groovy`

Showing health and status for app cf-scale-boot in org oreilly-class / space instructor as [mstine@pivotal.io](#)...

OK

requested state: started

instances: 1/1

usage: 512M x 1 instances

urls: cf-scale-boot-stockinged-rust.cfapps.io

last uploaded: Fri Feb 13 18:56:29 UTC 2015

state	since	cpu	memory	disk
-------	-------	-----	--------	------

#0 running 2015-02-13 04:01:50 PM 0.0% 389.1M of 512M 128.9M of 1G

And again, view the additional app.update event:

```
$ cf events cf-scale-boot
```

Getting events for app cf-scale-boot in org oreilly-class / space instructor as [mstine@pivotal.io](mailto:mstine@pivotal.io)...

time event actor description

2015-02-13T16:01:28.00-0600 audit.app.update [mstine@pivotal.io](mailto:mstine@pivotal.io) state: STARTED

2015-02-13T15:59:10.00-0600 audit.app.update [mstine@pivotal.io](mailto:mstine@pivotal.io) state: STOPPED

2015-02-13T15:18:33.00-0600 audit.app.update [mstine@pivotal.io](mailto:mstine@pivotal.io) instances: 1

2015-02-13T15:04:34.00-0600 audit.app.update [mstine@pivotal.io](mailto:mstine@pivotal.io) instances: 5

2015-02-13T12:56:35.00-0600 audit.app.update [mstine@pivotal.io](mailto:mstine@pivotal.io) state: STARTED

2015-02-13T12:56:26.00-0600 audit.app.update [mstine@pivotal.io](mailto:mstine@pivotal.io)

2015-02-13T12:56:26.00-0600 audit.app.map-route [mstine@pivotal.io](mailto:mstine@pivotal.io)

2015-02-13T12:56:24.00-0600 audit.app.create [mstine@pivotal.io](mailto:mstine@pivotal.io) instances: 1,  
memory: 512, state: STOPPED, environment\_json: PRIVATE DATA HIDDEN

Logs

One of the most important enablers of visibility into application behavior is logging. Effective management of logs has historically been very difficult. Cloud Foundry's log aggregation components simplify log management by assuming responsibility for it. Application developers need only log all messages to either STDOUT or STDERR, and the platform will capture these messages.

For Developers

Application developers can view application logs using the CF CLI.

Let's view recent log messages for cf-scale-boot:

```
$ cf logs cf-scale-boot --recent
```

Here are two interesting subsets of one output from that command:

Example 1. CF Component Logs

2015-02-13T14:45:39.40-0600 [RTR/0] OUT [cf-scale-boot-stockinged-rust.cfapps.io](http://cf-scale-boot-stockinged-rust.cfapps.io)

- [13/02/2015:20:45:39 +0000] "GET /css/bootstrap.min.css HTTP/1.1" 304 0  
"http://cf-scale-boot-stockinged-rust.cfapps.io/" "Mozilla/5.0 (Macintosh; Intel Mac OS X 10\_9\_5) AppleWebKit/537.36 (KHTML, like Gecko) Chrome/40.0.2214.111 Safari/537.36" 10.10.66.88:50372 x\_forwarded\_for:"50.157.39.197"  
vcap\_request\_id:84cc1b7a-bb30-4355-7512-5adaf36ff767  
response\_time:0.013115764 app\_id:7a428901-1691-4cce-b7f6-62d186c5cb55 (1)  
2015-02-13T14:45:39.40-0600 [RTR/1] OUT cf-scale-boot-stockinged-rust.cfapps.io  
- [13/02/2015:20:45:39 +0000] "GET /img/LOGO\_CloudFoundry\_Large.png HTTP/1.1" 304 0 "http://cf-scale-boot-stockinged-rust.cfapps.io/" "Mozilla/5.0 (Macintosh; Intel Mac OS X 10\_9\_5) AppleWebKit/537.36 (KHTML, like Gecko) Chrome/40.0.2214.111 Safari/537.36" 10.10.66.88:24323  
x\_forwarded\_for:"50.157.39.197" vcap\_request\_id:b3e2466b-6a41-4c6d-5b3d-0f70702c0ec1 response\_time:0.010003444 app\_id:7a428901-1691-4cce-b7f6-62d186c5cb55  
2015-02-13T15:04:33.09-0600 [API/1] OUT Tried to stop app that never received a start event (2)  
2015-02-13T15:04:33.51-0600 [DEA/12] OUT Starting app instance (index 2) with guid 7a428901-1691-4cce-b7f6-62d186c5cb55 (3)  
2015-02-13T15:04:33.71-0600 [DEA/4] OUT Starting app instance (index 3) with guid 7a428901-1691-4cce-b7f6-62d186c5cb55  
An "Apache-style" access log event from the (Go)Router

An API log event that corresponds to an event as shown in cf events

A DEA log event indicating the start of an application instance on that DEA.

## Example 2. Application Logs

2015-02-13T16:01:50.28-0600 [App/0] OUT 2015-02-13 22:01:50.282 INFO 36 — [runner-0] o.s.b.a.e.jmx.EndpointMBeanExporter : Located managed bean 'autoConfigurationAuditEndpoint': registering with JMX server as MBean [org.springframework.boot:type=Endpoint,name=autoConfigurationAuditEndpoint]  
2015-02-13T16:01:50.28-0600 [App/0] OUT 2015-02-13 22:01:50.287 INFO 36 — [runner-0] o.s.b.a.e.jmx.EndpointMBeanExporter : Located managed bean 'shutdownEndpoint': registering with JMX server as MBean [org.springframework.boot:type=Endpoint,name=shutdownEndpoint]  
2015-02-13T16:01:50.29-0600 [App/0] OUT 2015-02-13 22:01:50.299 INFO 36 — [runner-0] o.s.b.a.e.jmx.EndpointMBeanExporter : Located managed bean

```
'configurationPropertiesReportEndpoint': registering with JMX server as MBean  
[org.springframework.boot:type=Endpoint,name=configurationPropertiesReportEndpoint]  
2015-02-13T16:01:50.36-0600 [App/0] OUT 2015-02-13 22:01:50.359 INFO 36 — [  
runner-0] s.b.c.e.t.TomcatEmbeddedServletContainer : Tomcat started on port(s):  
61316/http  
2015-02-13T16:01:50.36-0600 [App/0] OUT Started...  
2015-02-13T16:01:50.36-0600 [App/0] OUT 2015-02-13 22:01:50.364 INFO 36 — [  
runner-0] o.s.boot.SpringApplication : Started application in 6.906 seconds (JVM  
running for 15.65)
```

As you can see, Cloud Foundry's log aggregation components capture both application logs and CF component logs relevant to your application. These events are properly interleaved based on time, giving you an accurate picture of events as they transpired across the system.

To get a running “tail” of the application logs rather than a dump, simply type:

```
$ cf logs cf-scale-boot
```

You can try various things like refreshing the browser and triggering stop/start events to see logs being generated.

## Health

---

1. Demonstrate scale demo application

# Lab 7 - Push Cloud Native Apps to Pivotal Cloud Foundry

---

## Build and Push!

---

1. Return to the Fortune Teller application directory that you finished in Lab 3:

```
$ cd $COURSE_HOME/labs/initial/fortune-teller
```



## 2. Using Maven, build and package the application:

```
$ mvn package
```

Maven will automatically download all of Fortune Teller's dependencies. This may take a few moments.

## 3. Push the Spring Cloud services:

```
$ cf push -f manifest-services.yml
```

This will push the Spring Cloud Config Server, a Eureka server, and a Hystrix Dashboard, all with random routes.

## 4. Edit scripts/create\_services.sh to add the random routes that were generated for you:

```
cf cups config-service -p '{"uri":"http://config-server-fluxional-sutt  
ee.cfapps.io"}'  
cf cups service-registry -p '{"uri":"http://eureka-unprevalent-toper.c  
fapps.io"}'  
cf cs elephantsql turtle fortunes-db
```

## 5. Run scripts/create-services.sh to create the services that you need:

```
$ scripts/create_services.sh
```

```
Creating user provided service config-service in org platform-eng / sp  
ace nfjs-workshop as mstine@pivotal.io...  
OK  
Creating user provided service service-registry in org platform-eng /  
space nfjs-workshop as mstine@pivotal.io...  
OK  
Creating service fortunes-db in org platform-eng / space nfjs-workshop  
as mstine@pivotal.io...  
OK
```

6. Push the microservices:

```
$ cf push -f manifest-apps.yml
```

This will push the fortunes service and the ui application.


7. In a browser, access the fortunes-ui application at the route that was created for you:

# Fortune Teller!

Find out what the future holds...

It is now, and in this world, that we must live.

8. Now, in another browser tab, access the Hystrix Dashboard at the route that was created for you. Enter the route for the UI application and click the “Monitor Stream.”



## Hystrix Dashboard

<http://fortunes-ui-polyunsaturated-hogg.cfapps.io/hystrix.stream>

*Cluster via Turbine (default cluster):* <http://turbine-hostname:port/turbine.stream>  
*Cluster via Turbine (custom cluster):* [http://turbine-hostname:port/turbine.stream?cluster=\[clusterName\]](http://turbine-hostname:port/turbine.stream?cluster=[clusterName])  
*Single Hystrix App:* <http://hystrix-app:port/hystrix.stream>


Delay:  ms

Title:


9. Access the fortunes-ui and show that the circuit breaker is registering successful requests.

**Hystrix Stream:** <http://fortunes-ui-polyunsaturated-hogg.cfapps.io/hystrix.stream>

**Circuit** Sort: [Error then Volume](#) | [Alphabetical](#) | [Volume](#) | [Error](#) | [Mean](#) | [Median](#) | [90](#) | [99](#) | [99.5](#)



**Thread Pools** Sort: [Alphabetical](#) | [Volume](#) |



10. Stop the fortunes application:

```
$ cf stop fortunes
```

11. Access the fortunes-ui and see that the “fallback fortune” is being returned.

# Fortune Teller!


Find out what the future holds...

Your future is unclear.


12. Access the fortunes-ui and show that the circuit breaker is registering short-circuited requests.

**Hystrix Stream:** <http://fortunes-ui-polyunsaturated-hogg.cfapps.io/hystrix.stream>

**Circuit** Sort: [Error then Volume](#) | [Alphabetical](#) | [Volume](#) | [Error](#) | [Mean](#) | [Median](#) | [90](#) | [99](#) | [99.5](#)



**Thread Pools** Sort: [Alphabetical](#) | [Volume](#) |



13. Start the fortunes application:

```
$ cf start fortunes
```

Continue to access the fortunes-ui and watch the dashboard. After the fortunes service has re-registered with Eureka and the fortunes-ui load balancer caches are refreshed, you will see the circuit breaker recover. You should then start getting random fortunes again!